

Format

DANFEVER contains three sqlite databases (SQLite Consortium, 2000); `da-fever.db`, `da-wikipedia.db` and `den-store-danske.db`.

`da-wikipedia.db` contains article data from Danish Wikipedia. It contains an `id`-field, which is the articles' `curid` from Wikipedia. It also contain the `text` and `title` of each article, as well as the `url` to that article.

`den-store-danske.db` contains the same information from Den Store Danske, but the article ID is a simple enumeration, and the article texts are not included due to data-right (evidence extracts are included though as described below).

The `da-fever.db` database contain the annotated claims. Each row in the database contain a `claim` and a unique `id`. With each claims comes the labels `verifiable` (`Verifiable` and `NotVerifiable`) and `label` (`Supported`, `Refuted` and `NotEnoughInfo`). The `evidence` column contain information about what articles were used to create and annotate the claim, and is composed by a comma-separated string, with IDs referring to the articles. The ID-format is `Y_X` where `Y` is either `wiki` or `dsd` to indicate whether the article comes from Danish Wikipedia or Den Store Danske, and `X` is the numerical id from that data-source. Finally the claims has a related `evidence_extract` that can be used to train inference models. Note that there may be some character-level incongruence between the original articles and the `evidence_extract`, due to formatting and scraping. The evidence extracts for `Verifiable` claims are the text-snippets used to create and annotate the claims. The evidence extract for `NotEnoughInformation` samples are based on a chrF1 search from the related article, in order to get relevant text.

All three databases are also provided in TSV-format.